

Master Thesis: Offer Networks: Simulation and Dynamics

Zar Goertzel

August 18, 2017

What is an Offer Network?

- In the literature: "Barter Exchange Networks".

What is an Offer Network?

In the literature: "Barter Exchange Networks".

Basic idea:

- Users exchange services and goods.
- Allow ($n > 2$)-user exchanges (also gifts and chains).
- Use optimization algorithms to suggest exchanges (matches).

What is an Offer Network?

In the literature: "Barter Exchange Networks".

Basic idea:

- Users exchange services and goods.
- Allow ($n > 2$)-user exchanges (also gifts and chains).
- Use optimization algorithms to suggest exchanges (matches).

But why? Isn't money more efficient?

What's wrong with money?

Answer 1: sometimes it's unethical.

- *Only one* country allows organ sale.
- Organ sale is outlawed. It's a *repugnant* transaction.

What's wrong with money?

Answer 1: sometimes it's unethical.

Only one country allows organ sale.

Organ sale is outlawed. It's a *repugnant* transaction.

- Result: Kidney Exchange Network boom!
 - First US kidney exchange: 2000.
 - Now: 140+ transplant center kidney exchange network with biweekly matching.

What's wrong with money?

Answer 2: National Odd Shoe Exchange.

- Otherwise you would pay double.

What's wrong with money?

Answer 3: Ambiguous value:

- Home exchange:
 - Market value doesn't have to match up perfectly to *subjectively* satisfy homeowners.

What's wrong with money?

Answer 3: Ambiguous value:

- Home exchange:
 - Market value doesn't have to match up perfectly to *subjectively* satisfy homeowners.
- Data exchange:
 - What is value of data anyway?
 - May want framework for carefully restricted access . . .

What's wrong with money?

Answer 4: Negligible value and scale:

- Book exchange.
- Homework help.
- Proofreading (of blogs or dating profiles – not just school essays).
- (Academic) book review exchange:
 - University professors would charge a high fee, unless the favor is returned.

++ I-YOU vs I-IT interactions

What's wrong with money?

Answer 5: Nothing.

- Works pretty well as market mechanism.
- Don't need "coincidence of wants".
 - Bookmooch uses credit system
- Recommender systems are still useful.
- *Exchange recommendation also useful?*
 - Elimination/reduction of middle-men.

State of Offer Network research

- Mostly kidney exchange research:

State of Offer Network research

Mostly kidney exchange research:

- 4 types of kidneys.
- Only interest in optimal solutions with small exchanges (≤ 3):
 - NP-Hard and APX-Complete.
 - Unless exchange-size unbounded with only edge-weight constraints.
- Usually Erdos-Renyi graph models

State of Offer Network research

Mostly kidney exchange research:

- Usually Erdos-Renyi graph models
- Domain-specific techniques used to improve Integer Linear Programming speed.

State of Offer Network research

Mostly kidney exchange research:

- Recently: PTIME approximation algorithm achieves 90 – 98%.
- Recently: Model incorporating rejection probability leads to 15%+ increase in transplants.
- Recently: Model keeping some kidneys for future use performs better (less *myopic*).
 - Latent potential for matches in graph is limited. Preserving useful parts helps.

State of Offer Network research

- Decentralized Offer Network: can selfishly exchanging agents pull off MacDonald's "paper clip for house" exchange chain?

State of Offer Network research

- Decentralized Offer Network: can selfishly exchanging agents pull off MacDonald's "paper clip for house" exchange chain? – Yes!

State of Offer Network research

- Decentralized Offer Network: can selfishly exchanging agents pull off MacDonald's "paper clip for house" exchange chain? – Yes!
- Greedy matching is best in Erdos-Renyi model.
 - Gift chains are better.

State of Offer Network research

- Abbassi test greedy cycle cover and two approximation algorithms on scale-free graph!
- Follow up with a comparison to credit-based system (works much better).
- Rappaz develop recommender system for barter exchange websites.

Question for this thesis

- 1 Extend Abbassi's experiments, and *include match acceptance probability*.

Question for this thesis

- 1 Extend Abbassi's experiments, and *include match acceptance probability*.
- 2 Is greedy matching frequency also best in scale-free graphs?
- 3 Is there a steady-state size of the Offer Network graph?

Question for this thesis

- 1 Extend Abbassi's experiments, and *include match acceptance probability*.
- 2 Is greedy matching frequency also best in scale-free graphs?
- 3 Is there a steady-state size of the Offer Network graph?
 - How long do users wait?
 - How many users are matched?
- 4 Does matching marginalize unpopular tasks?


Question for this thesis

- 1 Extend Abbassi's experiments, and *include match acceptance probability*.
- 2 Is greedy matching frequency also best in scale-free graphs?
- 3 Is there a steady-state size of the Offer Network graph?
 - How long do users wait?
 - How many users are matched?
- 4 Does matching marginalize unpopular tasks?
- 5 Alternative to rejecting whole match if one user rejects?

Question for this thesis

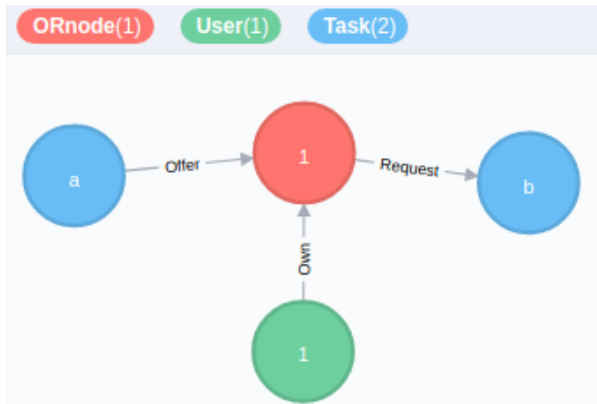
- Alternative to rejecting whole match if some users reject?

Very big exchange in "optimal matching" without size constraints:



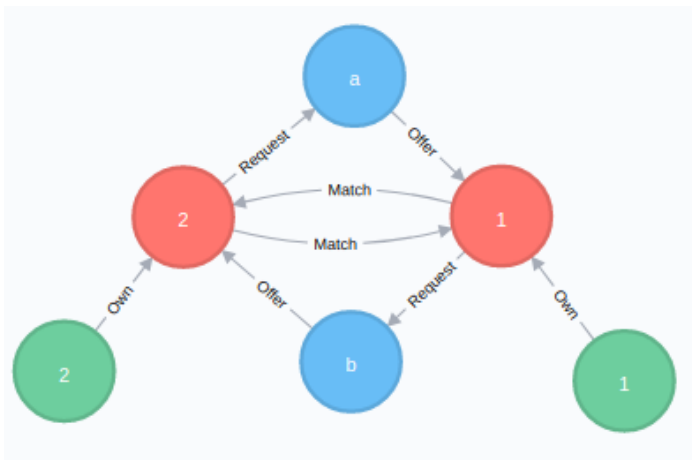
ORpair: (offer, request) pair

User 1 offers task a in exchange for requested task b:




ORpair: (offer, request) pair


A match between User 1 and User 2's ORpairs:



Scale-Free Graph (RB1) with 100 ORpairs



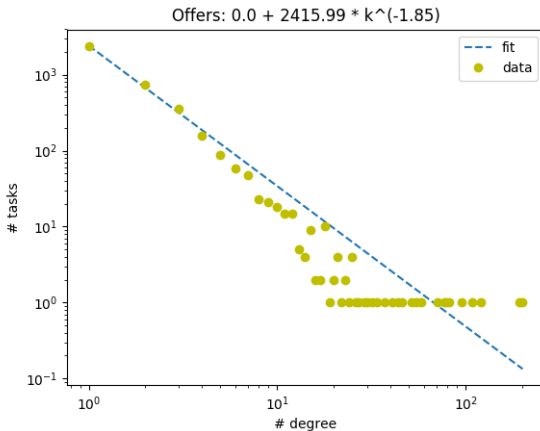
Erdos-Renyi ($p = 0.04$) with 100 ORpairs



Experimental Set-up: Data

- Find distribution of offers and requests in Ratebeer and Bookmooch datasets provided.
- Generate scale-free graphs with similar popularity distributions.
- I use 3 such graphs.

Scale-Free Graph (RB1) with 10,000 ORpairs



Experimental Set-up

- Time measured by new ORpairs
- Every n_{match} ORpairs, run a match algorithm
 - Assume uniform acceptance probability p .
- Add $N_{initial}$ ORpairs and run matching algorithm to initialize.
- Run test until N_{end} ORpairs added.

Experimental Set-up

- Time measured by new ORpairs
- Every n_{match} ORpairs, run a match algorithm
 - Assume uniform acceptance probability p .
- Add $N_{initial}$ ORpairs and run matching algorithm to initialize.
- Run test until N_{end} ORpairs added.
- Use same series of ORpair updates for each test run in an experiment.

Matching Algorithms

- Maximum 2-way matching: $\mathcal{O}(|2\text{-cycle}|^3)$. Performs horribly.
- Maximum Edge-Weight Matching (MAX). $\mathcal{O}(|ORpair|^3)$. Optimal if users always accept matches. I use Munkres algorithm implemented in Cython.
- Greedy Shortest Cycle (GSC). $\mathcal{O}((|ORpairs| + |tasks| \ln |tasks|)|ORpairs|)$. Good for $p < 1$?
- Dynamic Shortest Cycle (DYN). GSC only on new ORpairs.

Matching Algorithms

- Maximum 2-way matching: $\mathcal{O}(|2\text{-cycle}|^3)$. Performs horribly.
- Maximum Edge-Weight Matching (MAX). $\mathcal{O}(|ORpair|^3)$. Optimal if users always accept matches. Uses $|ORpair|^2$ matrix.
- Greedy Shortest Cycle (GSC).
 $\mathcal{O}((|ORpairs| + |tasks| \ln |tasks|) |ORpairs|)$. Good for $p < 1$?
 - Abbassi: use GSC many times.
 - Jia: use GSC to seed local search.
 - Jia: use product of degree (PoD) order. Bias toward unpopular tasks!

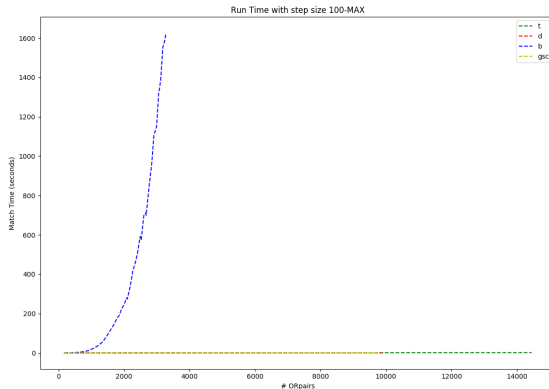
Neo4j Query Implementation of GSC

```
MATCH (o:ORnode)-[reqR:Request]->(req:Task),
p = shortestPath((req)-[link:Offer|:Request*]->(o))
WHERE NOT exists(reqR.matched) AND
ALL (r IN relationships(p) WHERE NOT exists(r.matched))
FOREACH (r IN link | SET r.matched = TRUE)
SET reqR.matched = TRUE
WITH FILTER(ornode IN nodes(p) WHERE ornode:ORnode) AS off
UNWIND p as off
MATCH (off)<-[]-()->[]-(req:ORnode)
WHERE req IN p AND off.offer = req.request
CREATE (off)-[:Match]->(req)
```

Run Time - MAX


Max run from 100 to 5000 ORpairs. Others 15,000.

$$n_{match} = 100.$$




Run Time

Run from 100 to 15,000 ORpairs. $n_{match} = 20$.



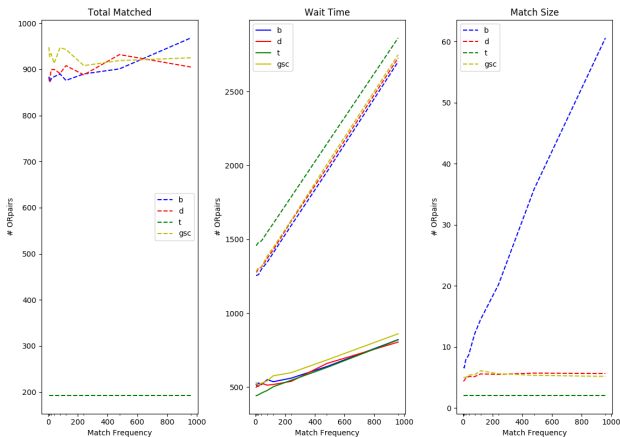
Run Time - GSC-PoD

Run from 100 to 15,000 ORpairs. $n_{match} = 100$.




Match Frequency Experiment: $p = 1$

Run from 100 to 3003 ORpairs.




$N_{initial}$ Experiment: $p = 1$

Run for 3000 steps.



Long Run Experiment: $p = 1$

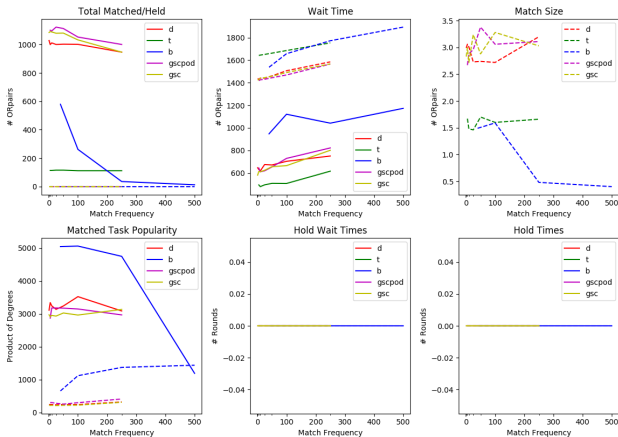
Run from 100 to 35,00 ORpairs. $n_{match} = 20$.



- Matching algorithms are all comparable for $p = 1$.
- Does MAX's performance deteriorate as expected?

Performance Test: $p = 0.9$

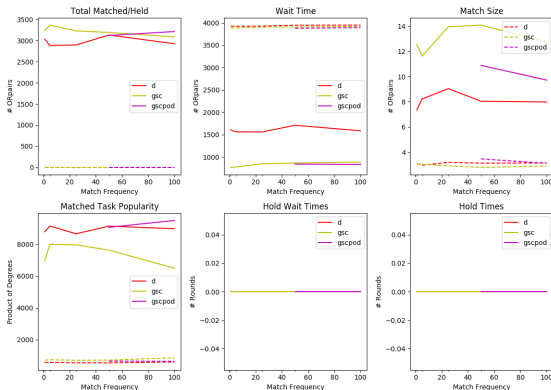
Run from 400 to 3400 (as MAX is slow)



- Matching algorithms are all comparable for $p = 1$.
- Does MAX's performance deteriorate as expected? – Yes.
- What about the other algorithms?
 - Wait time increases as p falls.
 - Accepted match size slowly falls.
 - Unpopular tasks marginalized. PoD less so.
 - Total matched ORpairs decreases drastically.**

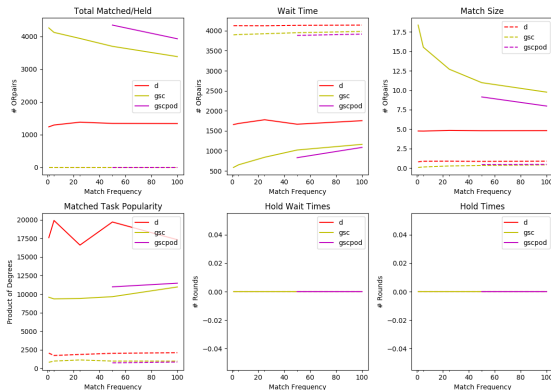
Performance Test: $p = 0.9$

Run from 10,000 to 15,000



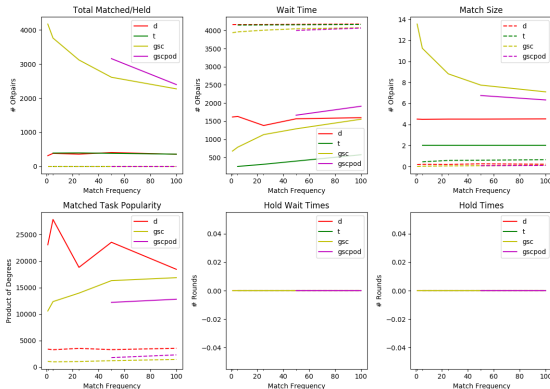
Performance Test: $p = 0.7$

Run from 10,000 to 15,000



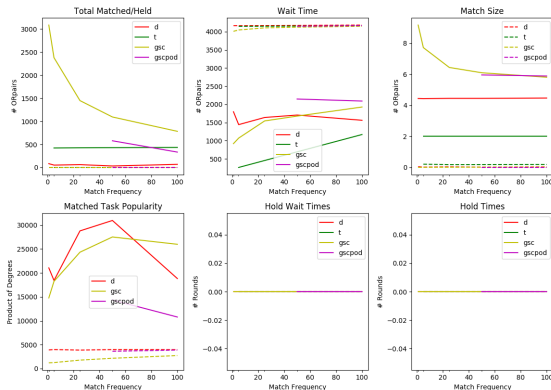
Performance Test: $p = 0.5$

Run from 10,000 to 15,000



Performance Test: $p = 0.3$

Run from 10,000 to 15,000



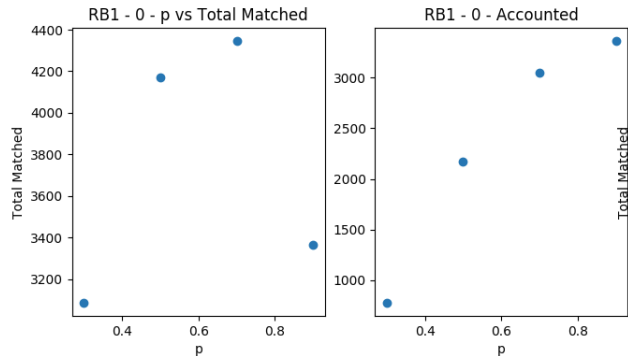
Two Sloppy/Mistaken Assumptions

- 1 During initialization, $N_{initial}$ nodes are added and then a matching algorithm is run once.
 - Problem: many more matched at $p = 0.9$ than $p = 0.3$. Thus there is more potential for matches left.
 - Measuring the number and accounting, the total matched nodes decrease with p .

2

Performance Test: p vs total matched

Run from 10,000 to 15,000



Two Sloppy/Mistaken Assumptions

- 1 During initialization, $N_{initial}$ nodes are added and then a matching algorithm is run once.
- 2 Rematching is allowed.
 - Lazy reasoning behind assumption: there will be better things to do than rematch.
 - Big problem: small step size and shortest cycle matching.
 - However, low acceptance probability cases *don't catch up*.


Can we maintain performance for low p ?

Hanging ORpair method:

- Assume rough parity among matched tasks.
- Only need to worry about ORpairs bordering a rejected ORpair: combine into one node and re-add to graph.

Held ORpair Example

If OR1 rejects, make new node (offer t2, request t1).



HOR: pros and cons

Pros:

- Algorithmically, ORpair acceptance rate is similar to that for 3-way-matches.
- Held ORpairs are semi-random, but resemble Dickerson's potentials approach?
I.e., less myopic
- Hint at generalization to asymmetric Offer Network design (or wait-list use). (*Future Work*)

HOR: pros and cons

Theoretical Cons:

- Parity doesn't always hold.
- Acceptance probability becomes p^2 ? (–Also forgot to implement.)
- Complicates additional features (e.g., task expiration, timing)

Practical Cons:

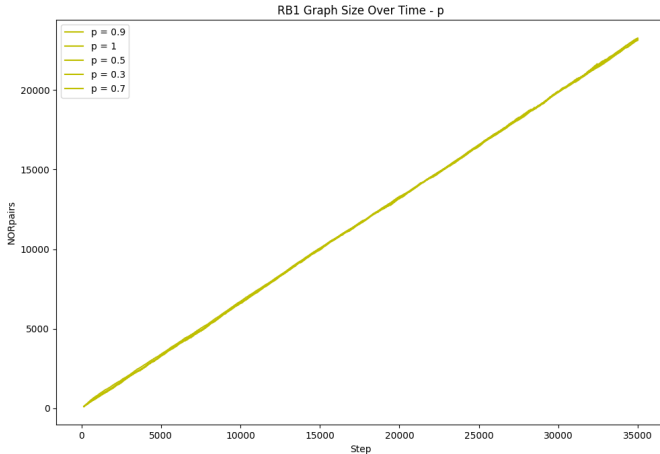
- User has to do more. Risks longer wait time.
- Incentive for user with held offer to stay?
 - Reputation system: negatively rate user for leaving: hard to find new matches.
 - Can be implemented like PoD or with MAX.
- Alternatively: add held request to waitlist queue; treat offer as gift.

HOR: Rematching

- 1 Held ORpair (t_2, t_1) can immediately be rematched with rejecting $(t_1, 2)$!
 - Murphy's law scenario: identical performance to $p = 1$ case with marginal wait time difference.
 - Curiously, potential ORpairs matched faster at lower p .
- 2 In the long run, latent graph potential for matches results in equal performance.
 - Properly disabling rematching may lead to low- p depleting graph potential.
 - However, this may merely mean postponed potential.

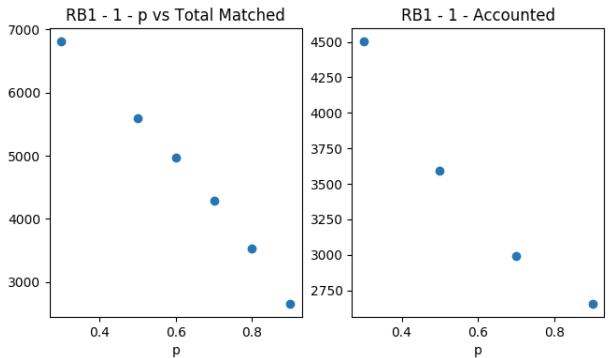
Long Run Graph Size

Run from 100 to 35,000




Performance Test: p vs total matched

Run from 10,000 to 15,000



Performance Test: p vs total matched

Ran a few runs with fixed, standardized ($p = 1$) matching for initialization.

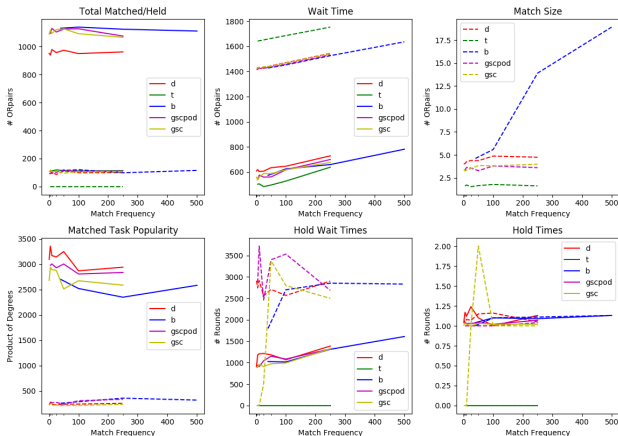


Performance Test: General Results

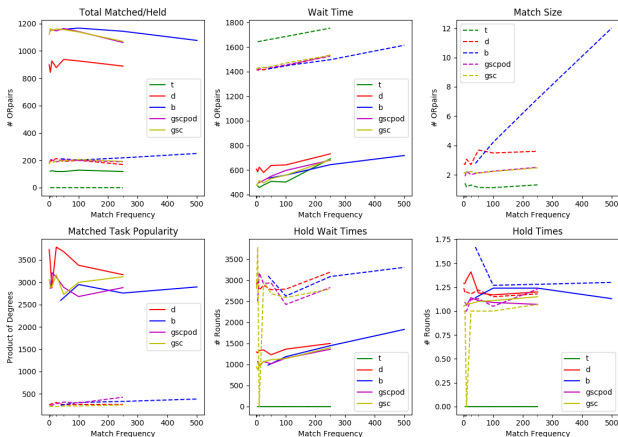
Tests with MAX: 400 to 3400

- With HOR, MAX performs best.
- Larger match size is okay.
- ORpairs held significantly longer
 - Which fits theoretical prediction. Perhaps abuses *rematching* less.
- Surprisingly: includes unpopular tasks best!

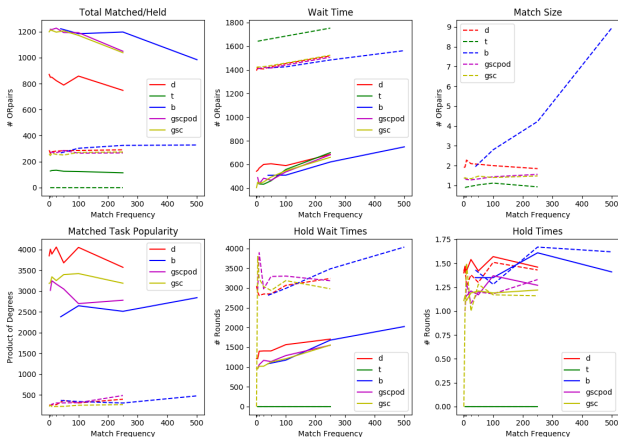
HOR Performance Test: $p = 0.9$



HOR Performance Test: $p = 0.8$

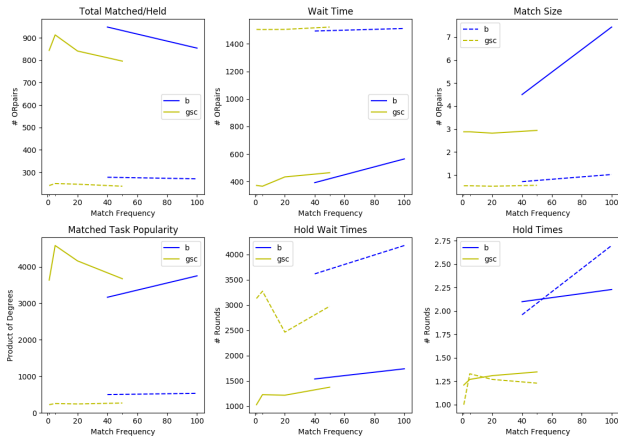


HOR Performance Test: $p = 0.7$



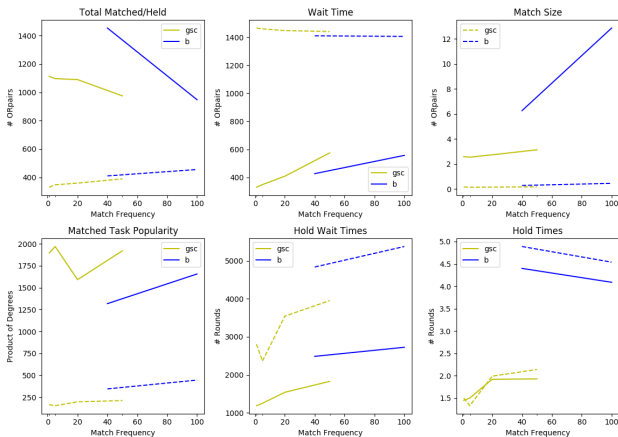
HOR Performance Test: $p = 0.5$

Note: using Bookmooch distribution graph



HOR Performance Test: $p = 0.3$

Note: using EN distribution graph.

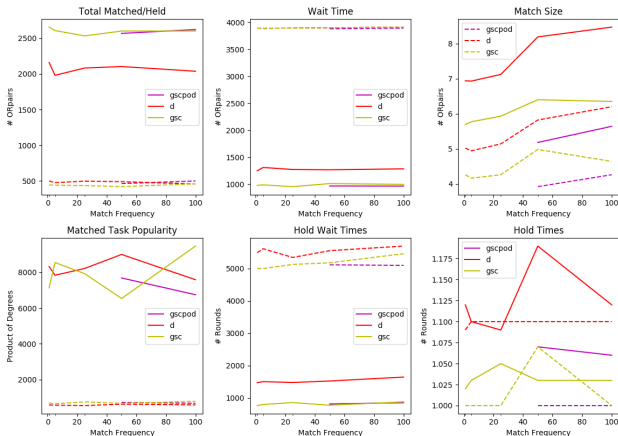


Performance Test: General Results

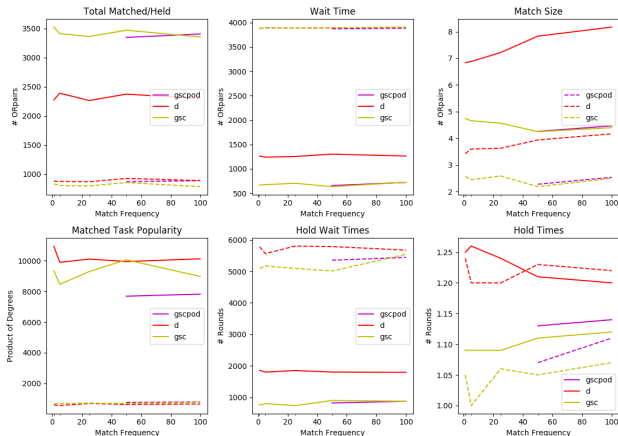
Tests from 10,000 to 15,000 (where wait time stabilizes)

- DYN does poorly below $p = 0.9$.
 - Because rejected/held ORpairs not re-added. (*future work*)
- GSC with Product of Degree order outperforms, marginally, in all metrics.
 - Q: is there a bigger difference without the ability to rematch ORpairs?
 - PoD much better for unpopular tasks at low p !
- Match size moves toward 2 and 3 as p decreases.
- Yet significantly more ORpairs are matched.
 - Postponed graph potential + shortening cycles?
 - A user's ORpair can be used in multiple conflicting cycles.

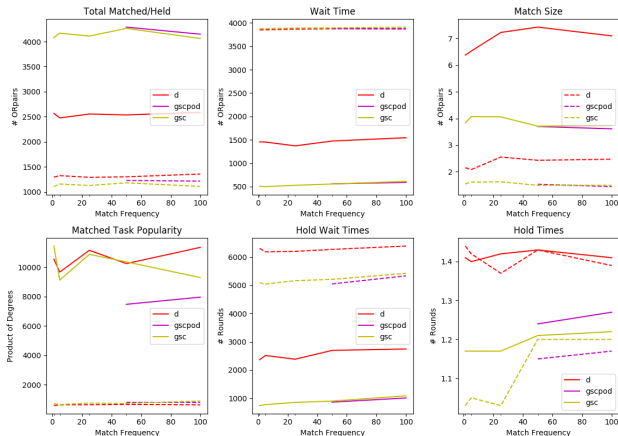
HOR Performance Test: $p = 0.9$



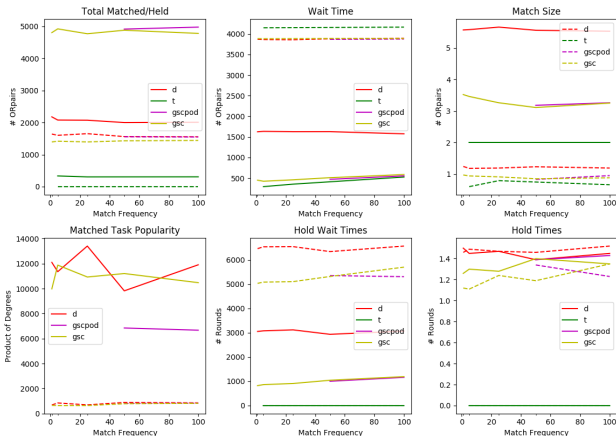
HOR Performance Test: $p = 0.8$



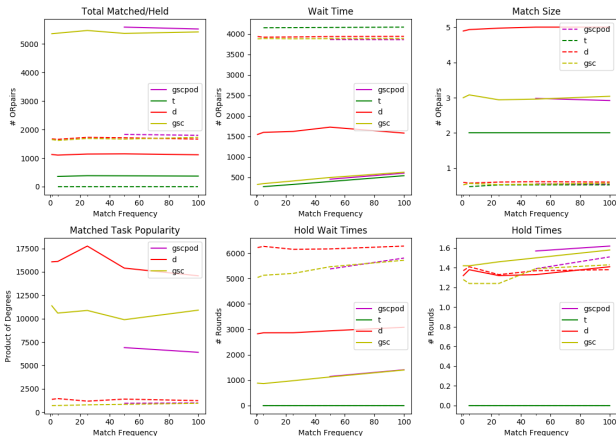
HOR Performance Test: $p = 0.7$



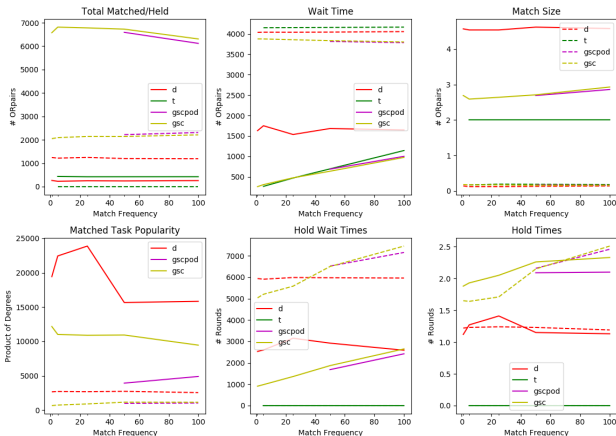
HOR Performance Test: $p = 0.6$



HOR Performance Test: $p = 0.5$



HOR Performance Test: $p = 0.3$



So, can Offer Networks work?

- 1 In the long run, roughly $1/3$ of ORpairs are matched.
 - Close to inherent graph limitation.
 - Task similarity based recommendation may help.

So, can Offer Networks work?

- 1 In the long run, roughly $1/3$ of ORpairs are matched.
 - Close to inherent graph limitation.
 - Task similarity based recommendation may help.
- 2 Wait-time appears to stabilize around 1000 steps.
 - If this holds for much larger graphs, realistic wait time may be ok.

So, can Offer Networks work?

- 1 In the long run, roughly $1/3$ of ORpairs are matched.
 - Close to inherent graph limitation.
 - Task similarity based recommendation may help.
- 2 Wait-time appears to stabilize around 1000 steps.
 - If this holds for much larger graphs, realistic wait time may be ok.
- 3 Matching Algorithms (beyond swaps) seem to matter more for speed of depleting graph of matches.
 - Good for use in decentralized case.

So, can Offer Networks work?

- 1 In the long run, roughly $1/3$ of ORpairs are matched.
 - Close to inherent graph limitation.
 - Task similarity based recommendation may help.
- 2 Wait-time appears to stabilize around 1000 steps.
 - If this holds for much larger graphs, realistic wait time may be ok.
- 3 Matching Algorithms (beyond swaps) seem to matter more for speed of depleting graph of matches.
 - Good for use in decentralized case.
- 4 Hold rounds get high (for MAX) with low- p . Could be cumbersome.
- 5 Low- p cases work best suggesting very many matches in high frequency

Future Work

- 1 Disallow rematching, and re-run experiments.
- 2 Impose match suggestion limits (irrespective of rematches).
- 3 Extend HOR to an asynchronous model (with gift chains allowed).
- 4 Experiment with task similarity and less uniform acceptance probability.
- 5 Add ORpair expiration.
- 6 Add reputation or user preferences to the equation.
- 7 Conjunctions and disjunctions to allow more intricate offers and requests.
(Will make the size of MAX's matrix bigger.)

Conclusions

- ➊ Given the framework, offer network style exchange recommendations can supplement traditional money marketplaces, and recommendation systems.
- ➋ But as in this thesis, probably not replace them.
(I don't have statistics on how long buyers and sellers wait on Amazon though.)
- ➌ The faster performance of low- p with HOR indicate that myopic matching isn't enough (as Dickerson found in the limited Kidney case.)
- ➍ As with combined organ exchanges (kidney and lung), linking specialized exchange markets should prove beneficial.

